

Michael Mihaley

mjm2442@columbia.edu | <https://github.com/michael-mih> | <https://www.linkedin.com/in/michael-mihaley-270626207/>

EDUCATION

Columbia University

Bachelor of Science in Computer Science, minor in Political Science

New York, NY

Sep. 2024 - May 2027

City College of New York (Macaulay Honors College)

Bachelor of Science in Computer Science, Dean's List

New York, NY

Sep. 2023 - May 2024 (transferred)

EXPERIENCE

Robotics and Machine Learning Researcher

May 2025 – Sep. 2025

Columbia University, Advisors: Yunzhu Li and Yixuan Wang

New York, NY

- Implemented Python/C++ control pipelines for 2 robotic systems (*Dexmate Vega*, *xArm7*), supporting joint-space and IK-based end-effector control with sub-cm positioning error
- Designed leader arm and extended the *GELLO Teleoperation Framework* to perform imitation reinforcement learning with the *Dexmate Vega*, simulating in *SAPIEN* and supporting 500 Hz control loops
- Designed and printed custom mounting connections to expand end effector options for the *Dexmate Vega* by 50%

Adjunct Faculty, Robotics

Oct. 2025 – Present

Rudolf Steiner School

New York, NY

- Instructing a *First Tech Challenge* robotics team across mechanical design, electronics, and Java programming, guiding them to achieve competition-ready results
- Introducing students to advanced programming concepts such as finite state machines and multi-threading to push the boundaries of code execution time
- Scaled program size from 2-3 students to around 15 in one semester through hands-on curriculum design

Machine Learning Intern

January 2024 – June 2024

Outlier AI

Remote

- Refined large language models via reference-free evaluation in Python for a response accuracy increase of over 40%
- Built automated safety checks combining classifiers and rule tests to flag 10% of pre-deployed output as misaligned

PROJECTS

Performance Optimized LLM Inference for Real-Time Robotic Planning | *Python, CUDA*

- Built a robotics planning pipeline using *TinyLlama* (1.1B) to convert natural language commands into action sequences, conducted full-stack GPU profiling (compute vs memory-bound analysis).
- Identified KV cache and attention layers as primary memory-bandwidth bottlenecks with benchmarking via *PyTorch Profiler*; reduced end-to-end latency by 40%, increased tokens/sec by 3×, and decreased VRAM usage by > 60% through 4-bit quantization and vLLM deployment.
- Evaluated latency-throughput tradeoffs under single-query (robotics) vs batched (cloud) workloads, characterizing memory scaling with sequence length for embedded deployment feasibility.
- Built end-to-end pipeline integrating LLM inference with *ROS2* and *Gazebo* simulation for performing generated action sequences.

HybVIO Visual-Inertial Odometry System with SLAM | *C++, C*

- Extended a research-oriented odometry system that uses preprocessed data to interface with actual *Intel RealSense* Cameras for real-time spatial mapping under constrained compute (Quad-core, 1.2GHz clock speed)
- Collaborated with a team to identify slowdowns in stereo imaging code and speed up visualization by 60%

CUDA SHA-3 Hash Checker | *CUDA, C++, C*

- Developed an NVIDIA CUDA parallelized SHA-3 hasher and checker, allowing the SHA-3 hashing and comparison of over 14,000,000 keywords in under 10 seconds
- Wrote a memory-mapped file processor to reduce wordlist processing time from 15 seconds to 0.5 seconds when working with millions of lines of input data

TECHNICAL SKILLS

Languages: C++, Java, CUDA, Python, JavaScript, Kotlin, Typescript, C#

Development Tools: Docker, Android Studio, Fusion 360, SAPIEN, Drake, CMake, MuJoCo, Pytorch, Linux, ROS

Certifications and Fellowships: CAIAC AI Policy Fellow, Coursera Deep Learning Certified